

CS2106 Introduction to Operating Systems

Semester 2 2023/2024

Questions

File Abstraction

1. Explain the following concepts in your own words clearly; the shorter, the better! Your explanation should be easily understandable for non-CS2106 students.

- What is a file?
- Name and describe some classifications of files.
- Distinguish between a file type and a file extension.
- What does it mean to open and close a file?
- What does it mean to truncate a file?

2. [Understanding directory permission] In Unix system, a directory has the same set of permission settings as a file. For example:

```
sooyj@sunfire [13:22:52] ~/tmp/Parent $ ls -l
total 8
drwx--x--x  2 sooyj  compsc   4096 Nov  8 13:22 Directory
sooyj@sunfire [13:22:53] ~/tmp/Parent $
```

You can see that directory **Directory** has the read, write, execute permission for owner, but only execution permission for group and others. It is easy to understand the permission bits for a regular file (read = can only access, write = can modify, execute = can execute this file). However, the same cannot be said for the directory permission bits.

Let's perform a few experiments to understand the permission bits for a directory.

Setup:

- Unzip **DirExp.zip** on any unix platform (Solaris, Mac OS X included).
- Change directory to the **DirExp/** directory, there are 4 subdirectories with the same set of files. Let's set their permission as follows:

| | |
|-----------------------|---|
| chmod 700 NormDir | NormDir is a normal directory with read, write and execute permissions. |
| chmod 500 ReadExeDir | ReadExeDir has read and execute permission. |
| chmod 300 WriteExeDir | WriteExeDir has write and execute permission. |
| chmod 100 ExeOnlyDir | ExeOnlyDir has only execute permission. |

Perform the following operations on each of the directory and note down the result. Make sure you are at the **DirExp/** directory at the beginning. DDDD is one of the subdirectories.

- a. Perform `"ls -l DDDD"`.
- b. Change into the directory using `"cd DDDD"`.
- c. Perform `"ls -l"`.
- d. Perform `"cat file.txt"` to read the file content.
- e. Perform `"touch file.txt"` to modify the file.
- f. Perform `"touch newfile.txt"` to create a new file.

Can you deduce the meaning of the permission bits for directory after the above? Can you use the "directory entry" idea to explain the behavior?

3. [Wrapping File Operations] File operations are very expensive in terms of time. There are several reasons: a) As we learned in the lecture, each file operation is a system call, which requires an execution mode change (user → kernel); b) Secondary storage mediums have high access latencies.

This leads to a strange phenomenon: it is generally true that the total time to perform 100 file operations for 1 item each is **much longer** than performing a single file operation for 100 items instead. e.g. writing one byte 100 times takes longer than writing 100 bytes in one go.

Most high-level programming languages therefore provide **buffered file operations** that wrap around primitive file operations. The buffered version maintains an internal intermediate storage in memory (i.e. buffer) to store values read from/written to the file by the user. For example, a **buffered write operation** will wait until the internal memory buffer is full before doing a large one-time file write operation to *flush* the buffer content into file.

- a. (Generalization) Give one or two examples of buffered file operations found in your favorite programming language(s). Other than the "chunky" read/write benefit, are there any other additional features provided by these high-level buffered file operations?
- b. (Application) Take a look at the given `"weird.c"` source code. Compile and perform the following experiments: Change the trigger value from 100, 200, ... until you see values printed on screen **before the program crashes**. Can you explain both the behavior and the significance of the "trigger" value? If you add a new line character `"\n"` to the `printf()` statement, how does the output pattern change? How can this information be useful?
- c. (Design) Give an algorithm in **high-level pseudo-code** to provide a buffered read operation. Use the following function header as a starting point:

```
BufferedFileRead(file, outputArray, arraySize)
// Read arraySize items from file and place the items in outputArray
```

4. [Wrapping File Operations, again] Study the attached program `weird_read.c`. Note that it is written to run with the given input `alice.txt`, but you can modify it to read your own file.
 - a. Uncomment `a()`; in `main`, then compile and run the program. What do you observe and why?
 - b. Uncomment `b()`; in `main`, then compile and run the program. What do you observe and why?
 - c. Uncomment `c()`; in `main`, then compile and run the program. What do you observe and why?
 - i. What happens when the first `read` and `printf` are removed instead?
 - ii. What happens when `c()` is run with multiple threads instead of processes?
-